

Classificação de Patologias usando Imagens Médicas

Carregar imagens do diretório

```
In [1]: import os
current_dir = os.path.abspath(os.getcwd())
```

Converter base de dados para treino, validação e teste

```
In [2]: #cria nova pasta para cachorros e gatos atendendo a estrutura do Keras/Tensorflow
folder = "/novo"
train_folder = current_dir + folder + "/train"
val_folder = current_dir + folder + "/val"
test_folder = current_dir + folder + "/test"

model_filepath = "keras/classificacao_02_03.keras"
conversao_path = "conversao/conversao_02_03"
```

Fazer o Tensorflow carregar as imagens para a RNA

```
In [3]: import tensorflow as tf

print(tf.config.list_physical_devices('GPU'))
print(tf.__version__)
```

```
2022-06-29 07:46:47.347256: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot open shared object file: No such file or directory
2022-06-29 07:46:47.347272: I tensorflow/stream_executor/cuda/cudart_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.
```

```
[ ]
2.6.1
```

```
2022-06-29 07:46:47.995773: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot open shared object file: No such file or directory
2022-06-29 07:46:47.995789: W tensorflow/stream_executor/cuda/cuda_driver.cc:269] failed call to cuInit: UNKNOWN ERROR (303)
2022-06-29 07:46:47.995800: I tensorflow/stream_executor/cuda/cuda_diagnostic_utils.cc:156] kernel driver does not appear to be running on this host (pc): /proc/driver/nvidia/version does not exist
```

```
In [4]: from tensorflow.keras.utils import image_dataset_from_directory
#image_dataset_from_directory monta uma estrutura de dados com imagens 180x180
# de 32 em 32 imagens
train_dataset = image_dataset_from_directory(train_folder, image_size=(180, 180),
validation_dataset = image_dataset_from_directory(val_folder, image_size=(180, 180),
```

```
test_dataset = image_dataset_from_directory(test_folder, image_size=(180, 180
```

Found 34931 files belonging to 2 classes.

Found 16 files belonging to 2 classes.

Found 484 files belonging to 2 classes.

2022-06-29 07:46:48.493037: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

In [5]:

```
#
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    print(data_batch[0].shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

(180, 180, 3)

2022-06-29 07:46:48.543356: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)

Treinando o modelo

In [6]:

```
from tensorflow import keras
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.layers.experimental.preprocessing import Rescaling

#cria uma arquitetura de uma rede neural profunda vazia
model = keras.Sequential()
model.add(Rescaling(scale=1.0/255))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(1, activation="sigmoid"))
model.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accuracy"])
#model.add(Dense(4, activation='softmax'))
#model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['ac
```

In [7]:

```
from tensorflow.keras.callbacks import ModelCheckpoint

callbacks = [
    ModelCheckpoint(
        filepath = model_filepath,
        save_best_only = True,
        monitor = "val_loss"
    )
]

history = model.fit(
    train_dataset,
    epochs=50,
```

```
validation_data=validation_dataset,  
callbacks=callbacks)
```

```
Epoch 1/50  
1092/1092 [=====] - 350s 320ms/step - loss: 0.4272 -  
accuracy: 0.8083 - val_loss: 0.1753 - val_accuracy: 0.9375  
Epoch 2/50  
1092/1092 [=====] - 342s 313ms/step - loss: 0.2633 -  
accuracy: 0.8977 - val_loss: 0.1099 - val_accuracy: 1.0000  
Epoch 3/50  
1092/1092 [=====] - 341s 312ms/step - loss: 0.2064 -  
accuracy: 0.9256 - val_loss: 0.0315 - val_accuracy: 1.0000  
Epoch 4/50  
1092/1092 [=====] - 349s 319ms/step - loss: 0.1465 -  
accuracy: 0.9501 - val_loss: 0.0277 - val_accuracy: 1.0000  
Epoch 5/50  
1092/1092 [=====] - 343s 314ms/step - loss: 0.0887 -  
accuracy: 0.9719 - val_loss: 0.0146 - val_accuracy: 1.0000  
Epoch 6/50  
1092/1092 [=====] - 342s 313ms/step - loss: 0.0493 -  
accuracy: 0.9853 - val_loss: 0.0093 - val_accuracy: 1.0000  
Epoch 7/50  
1092/1092 [=====] - 342s 313ms/step - loss: 0.0257 -  
accuracy: 0.9931 - val_loss: 0.0085 - val_accuracy: 1.0000  
Epoch 8/50  
1092/1092 [=====] - 345s 316ms/step - loss: 0.0166 -  
accuracy: 0.9962 - val_loss: 0.0079 - val_accuracy: 1.0000  
Epoch 9/50  
1092/1092 [=====] - 345s 316ms/step - loss: 0.0133 -  
accuracy: 0.9965 - val_loss: 0.0039 - val_accuracy: 1.0000  
Epoch 10/50  
1092/1092 [=====] - 341s 312ms/step - loss: 0.0117 -  
accuracy: 0.9968 - val_loss: 0.1475 - val_accuracy: 0.9375  
Epoch 11/50  
1092/1092 [=====] - 346s 316ms/step - loss: 0.0094 -  
accuracy: 0.9975 - val_loss: 0.0018 - val_accuracy: 1.0000  
Epoch 12/50  
1092/1092 [=====] - 356s 326ms/step - loss: 0.0086 -  
accuracy: 0.9975 - val_loss: 0.1214 - val_accuracy: 0.9375  
Epoch 13/50  
1092/1092 [=====] - 361s 330ms/step - loss: 0.0069 -  
accuracy: 0.9984 - val_loss: 0.3659 - val_accuracy: 0.9375  
Epoch 14/50  
1092/1092 [=====] - 346s 316ms/step - loss: 0.0083 -  
accuracy: 0.9975 - val_loss: 0.4307 - val_accuracy: 0.9375  
Epoch 15/50  
1092/1092 [=====] - 430s 394ms/step - loss: 0.0066 -  
accuracy: 0.9979 - val_loss: 0.0425 - val_accuracy: 1.0000  
Epoch 16/50  
1092/1092 [=====] - 408s 374ms/step - loss: 0.0072 -  
accuracy: 0.9985 - val_loss: 6.5603e-04 - val_accuracy: 1.0000  
Epoch 17/50  
1092/1092 [=====] - 392s 358ms/step - loss: 0.0065 -  
accuracy: 0.9988 - val_loss: 6.8530e-04 - val_accuracy: 1.0000  
Epoch 18/50  
1092/1092 [=====] - 416s 381ms/step - loss: 0.0073 -  
accuracy: 0.9985 - val_loss: 0.0065 - val_accuracy: 1.0000  
Epoch 19/50  
1092/1092 [=====] - 376s 344ms/step - loss: 0.0056 -  
accuracy: 0.9987 - val_loss: 0.0044 - val_accuracy: 1.0000  
Epoch 20/50  
1092/1092 [=====] - 343s 314ms/step - loss: 0.0038 -  
accuracy: 0.9993 - val_loss: 0.0039 - val_accuracy: 1.0000  
Epoch 21/50
```

```
1092/1092 [=====] - 352s 322ms/step - loss: 0.0053 -  
accuracy: 0.9989 - val_loss: 8.0096e-05 - val_accuracy: 1.0000  
Epoch 22/50  
1092/1092 [=====] - 357s 326ms/step - loss: 0.0054 -  
accuracy: 0.9987 - val_loss: 2.4788e-04 - val_accuracy: 1.0000  
Epoch 23/50  
1092/1092 [=====] - 353s 323ms/step - loss: 0.0079 -  
accuracy: 0.9981 - val_loss: 3.5050e-05 - val_accuracy: 1.0000  
Epoch 24/50  
1092/1092 [=====] - 406s 372ms/step - loss: 0.0049 -  
accuracy: 0.9991 - val_loss: 0.0109 - val_accuracy: 1.0000  
Epoch 25/50  
1092/1092 [=====] - 328s 301ms/step - loss: 0.0035 -  
accuracy: 0.9993 - val_loss: 2.7066e-05 - val_accuracy: 1.0000  
Epoch 26/50  
1092/1092 [=====] - 384s 352ms/step - loss: 0.0058 -  
accuracy: 0.9990 - val_loss: 7.4073e-06 - val_accuracy: 1.0000  
Epoch 27/50  
1092/1092 [=====] - 418s 382ms/step - loss: 0.0039 -  
accuracy: 0.9989 - val_loss: 0.0182 - val_accuracy: 1.0000  
Epoch 28/50  
1092/1092 [=====] - 355s 324ms/step - loss: 0.0038 -  
accuracy: 0.9993 - val_loss: 0.0106 - val_accuracy: 1.0000  
Epoch 29/50  
1092/1092 [=====] - 361s 330ms/step - loss: 0.0031 -  
accuracy: 0.9994 - val_loss: 0.0044 - val_accuracy: 1.0000  
Epoch 30/50  
1092/1092 [=====] - 358s 328ms/step - loss: 0.0068 -  
accuracy: 0.9986 - val_loss: 0.0713 - val_accuracy: 0.9375  
Epoch 31/50  
1092/1092 [=====] - 335s 307ms/step - loss: 0.0039 -  
accuracy: 0.9993 - val_loss: 3.5125e-06 - val_accuracy: 1.0000  
Epoch 32/50  
1092/1092 [=====] - 343s 314ms/step - loss: 0.0048 -  
accuracy: 0.9990 - val_loss: 3.6036e-04 - val_accuracy: 1.0000  
Epoch 33/50  
1092/1092 [=====] - 348s 319ms/step - loss: 0.0033 -  
accuracy: 0.9995 - val_loss: 0.0012 - val_accuracy: 1.0000  
Epoch 34/50  
1092/1092 [=====] - 350s 321ms/step - loss: 0.0051 -  
accuracy: 0.9987 - val_loss: 8.2792e-07 - val_accuracy: 1.0000  
Epoch 35/50  
1092/1092 [=====] - 334s 305ms/step - loss: 0.0029 -  
accuracy: 0.9995 - val_loss: 2.8217e-05 - val_accuracy: 1.0000  
Epoch 36/50  
1092/1092 [=====] - 338s 309ms/step - loss: 0.0022 -  
accuracy: 0.9995 - val_loss: 7.7823e-05 - val_accuracy: 1.0000  
Epoch 37/50  
1092/1092 [=====] - 342s 313ms/step - loss: 0.0062 -  
accuracy: 0.9987 - val_loss: 0.2868 - val_accuracy: 0.9375  
Epoch 38/50  
1092/1092 [=====] - 333s 305ms/step - loss: 0.0038 -  
accuracy: 0.9993 - val_loss: 0.3457 - val_accuracy: 0.9375  
Epoch 39/50  
1092/1092 [=====] - 334s 305ms/step - loss: 0.0028 -  
accuracy: 0.9995 - val_loss: 0.0171 - val_accuracy: 1.0000  
Epoch 40/50  
1092/1092 [=====] - 332s 304ms/step - loss: 0.0035 -  
accuracy: 0.9993 - val_loss: 0.1829 - val_accuracy: 0.9375  
Epoch 41/50  
1092/1092 [=====] - 333s 305ms/step - loss: 0.0033 -  
accuracy: 0.9993 - val_loss: 0.0916 - val_accuracy: 0.9375  
Epoch 42/50  
1092/1092 [=====] - 327s 299ms/step - loss: 0.0043 -
```

```

accuracy: 0.9993 - val_loss: 0.0384 - val_accuracy: 1.0000
Epoch 43/50
1092/1092 [=====] - 326s 298ms/step - loss: 0.0037 -
accuracy: 0.9994 - val_loss: 0.0037 - val_accuracy: 1.0000
Epoch 44/50
1092/1092 [=====] - 325s 297ms/step - loss: 0.0045 -
accuracy: 0.9991 - val_loss: 0.0012 - val_accuracy: 1.0000
Epoch 45/50
1092/1092 [=====] - 321s 294ms/step - loss: 0.0031 -
accuracy: 0.9995 - val_loss: 5.7260e-05 - val_accuracy: 1.0000
Epoch 46/50
1092/1092 [=====] - 326s 298ms/step - loss: 0.0050 -
accuracy: 0.9993 - val_loss: 1.8811e-04 - val_accuracy: 1.0000
Epoch 47/50
1092/1092 [=====] - 325s 297ms/step - loss: 0.0032 -
accuracy: 0.9995 - val_loss: 2.4629e-06 - val_accuracy: 1.0000
Epoch 48/50
1092/1092 [=====] - 326s 298ms/step - loss: 0.0033 -
accuracy: 0.9994 - val_loss: 4.0221e-05 - val_accuracy: 1.0000
Epoch 49/50
1092/1092 [=====] - 325s 298ms/step - loss: 0.0034 -
accuracy: 0.9995 - val_loss: 1.0507e-04 - val_accuracy: 1.0000
Epoch 50/50
1092/1092 [=====] - 332s 304ms/step - loss: 0.0026 -
accuracy: 0.9996 - val_loss: 2.6431e-04 - val_accuracy: 1.0000

```

In [8]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
flatten (Flatten)	(None, 484416)	0
dense (Dense)	(None, 1)	484417
Total params: 503,809		
Trainable params: 503,809		
Non-trainable params: 0		

In [9]:

```

#https://www.tensorflow.org/js/tutorials/conversion/import_keras?hl=pt-br#ali
import tensorflowjs as tfjs
tfjs.converters.save_keras_model(model, conversao_path)

```

Visualização de Resultados

In [10]:

```

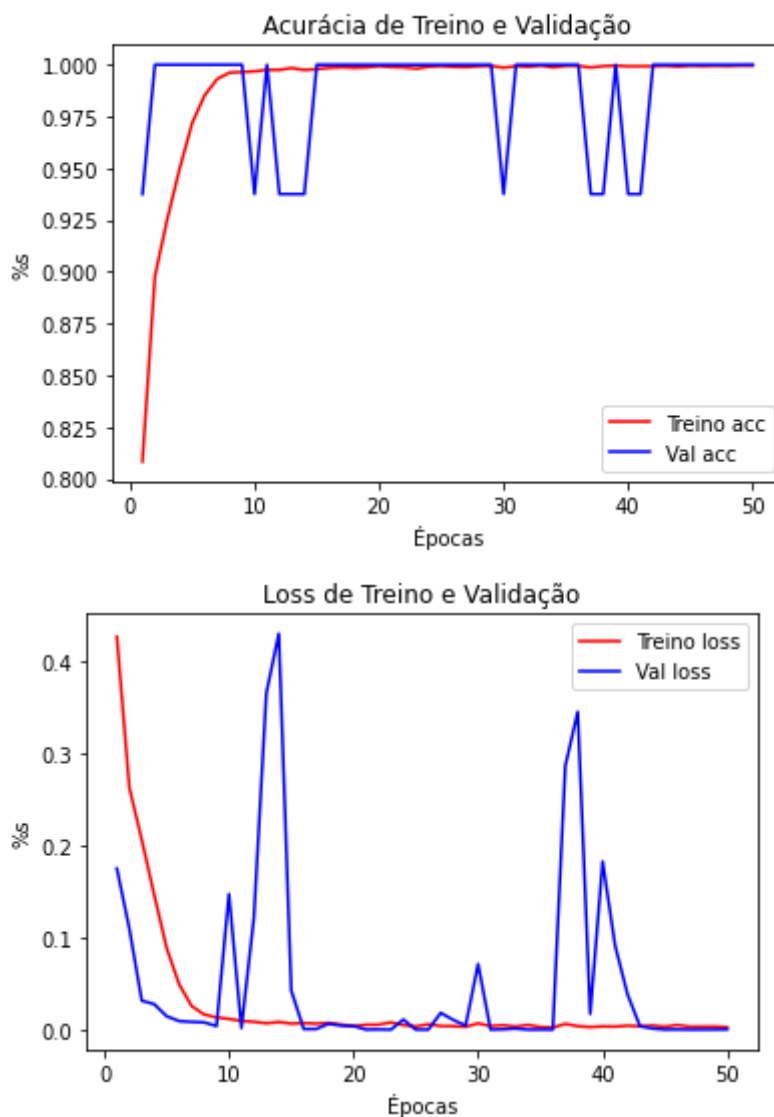
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]

```

```

epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "r", label="Treino acc")
plt.plot(epochs, val_accuracy, "b", label="Val acc")
plt.xlabel("Épocas")
plt.ylabel("%s")
plt.title("Acurácia de Treino e Validação")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "r", label="Treino loss")
plt.plot(epochs, val_loss, "b", label="Val loss")
plt.xlabel("Épocas")
plt.ylabel("%s")
plt.title("Loss de Treino e Validação")
plt.legend()
plt.show()

```



Resultados do Conjunto de Teste

```

In [11]: from tensorflow import keras
         model = keras.models.load_model(model_filepath)

```

```

In [12]: test_loss, test_acc = model.evaluate(test_dataset)
         print(f"Test accuracy: {test_acc:.3f}")

```

16/16 [=====] - 4s 71ms/step - loss: 0.0034 - accuracy: 1.0000
Test accuracy: 1.000

Referências

- <https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/>
- <https://stackoverflow.com/questions/3430372/how-do-i-get-the-full-path-of-the-current-files-directory>
- <https://www.geeksforgeeks.org/python-list-files-in-a-directory/>
- <https://pynative.com/python-random-sample/>
- <https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/>
- <https://www.mygreatlearning.com/blog/keras-tutorial/>
- <https://www.machinecurve.com/index.php/2020/03/30/how-to-use-conv2d-with-keras/>
- <https://www.pyimagesearch.com/2021/06/30/how-to-use-the-modelcheckpoint-callback-with-keras-and-tensorflow/>